

Linear Algebra 3

1. Let $A \in R^{m \times n}$ with rank = n; prove $\|A(A^T A)^{-1} A^T\|_2 = 1$

Since A is m x n a rectangular matrix with a rank n, it is full column rank

This implies that $m \geq n$

Since A is full column rank, this implies there exists a pseudoinverse of A

$$A^\dagger = (A^T A)^{-1} A^T$$

$$\|AA^\dagger\|_2 = 1$$

$$\|A(A^T A)^{-1} A^T\|_2 = 1$$

$$\|AA^{-1} A^{T^{-1}} A^T\|_2 = 1$$

$$\|IA^{T^{-1}} A^T\|_2 = 1$$

$$\|II\|_2 = 1 \rightarrow \|I\|_2 = 1$$

$$I = \begin{bmatrix} 1 & \dots & 0 \\ \dots & \dots & \dots \\ 0 & \dots & 1 \end{bmatrix}$$

Since you're talking the 2 – norm of the identity matrix, the max will always be 1

2. Let A and B be 2 semi-definite positive matrices in $R^{n \times n}$; Prove or disprove the following
- a. $A + B$ is positive semi-definite

If $A + B$ is positive semidefinite:

$$\forall x (x \in R^n \rightarrow x^T (A + B)x \geq 0)$$

Basically, for all x vectors in R^n , $x^T (A + B)x$ will always be greater than 0

$$x^T (A + B)x = x^T Ax + x^T Bx$$

Since A and B are positive semidefinite:

$$x^T Ax \geq 0 \text{ and } x^T Bx \geq 0$$

$$\text{Thus, } x^T (A + B)x = x^T Ax + x^T Bx \geq 0$$

TRUE

- b. AB is positive semi-definite

If AB is positive semidefinite:

$$\forall x (x \in R^n \rightarrow x^T (AB)x \geq 0)$$

Basically, for all x vectors in R^n , $x^T (AB)x$ will always be greater than 0

While we know $x^T Ax \geq 0$ and $x^T Bx \geq 0$

We don't know that $x^T ABx \geq 0$

FALSE

- c. B^T is positive semi-definite

If B^T is positive semidefinite:

$$\forall x (x \in R^n \rightarrow x^T (B^T)x \geq 0)$$

Basically, for all x vectors in R^n , $x^T (B^T)x$ will always be greater than 0

Since B is positive semidefinite, B is symmetric

Thus, $B^T = B$

$$x^T (B^T)x = x^T (B)x$$

Since we know $x^T Ax \geq 0$ and $x^T Bx \geq 0$, B^T is positive semidefinite

TRUE

Linear Classification

- Given a set of data points, $\{x^n\}$, we can define the convex hull to be the set of points given by $x = \sum_n \alpha_n x_n$, where $\alpha_n \geq 0$ and $\sum_n \alpha_n = 1$. Consider a second set of points, $\{y^n\}$, together, with their corresponding convex hull. By definition, the 2 sets of points will be linearly separable if there exists a vector, \hat{w} , and scalar, w_0 , such that $\hat{w}x_n + w_0 > 0$ for all x_n , and $\hat{w}y_n + w_0 < 0$ for all y_n . Show that if the 2 convex hulls intersect, their points are not linearly separable. Also show that if the 2 convex hulls don't intersect, they are linearly separable.

$$\hat{w}x_n + w_0 = \hat{w}\left(\sum_n \alpha_n x_n\right) + w_0$$

$$\hat{w}y_n + w_0 = \hat{w}\left(\sum_n \alpha_n y_n\right) + w_0$$

If the two functions intersect at some point:

$$\hat{w}^T x_n + w_0 \leq 0 \text{ and } \hat{w}^T y_m + w_0 \geq 0 \text{ at some point(s)}$$

$$\hat{w}^T x_n + w_0 \leq \hat{w}^T y_m + w_0 \rightarrow \hat{w}^T x_n \leq \hat{w}^T y_m$$

which means there is an x value that exists inside the convex hull of y ,

thus, they are not linearly separable; conversely, if no such points exist where this

is true, it means that there is a clear separation where all points, x_n , follow

$\hat{w}x_n + w_0 > 0$ and all points, y_n , follow $\hat{w}y_n + w_0 < 0$

2. By making use of $y = w^T x$, $m_k = w^T m_k$, and $s_k^2 = \sum_{n \in c_k} (y_n - m_k)^2$, show that the Fischer

Criterion: $J(w) = \frac{(m_2 - m_1)^2}{s_1^2 + s_2^2}$ can be written in the form: $J(w) = \frac{w^T S_B w}{w^T S_w w}$

$$S_B = (m_2 - m_1)(m_2 - m_1)^T$$

$$S_w = \sum_{n \in c_1} (x_n - m_1)(x_n - m_1)^T + \sum_{n \in c_2} (x_n - m_2)(x_n - m_2)^T$$

$$J(w) = \frac{(m_2 - m_1)^2}{s_1^2 + s_2^2}$$

$$J(w) = \frac{(w^T m_2 - w^T m_1)^2}{s_1^2 + s_2^2}$$

$$J(w) = \frac{(w^T m_2 - w^T m_1)^2}{\sum_{n \in c_1} (y_n - m_1)^2 + \sum_{n \in c_2} (y_n - m_2)^2}$$

$$J(w) = \frac{(w^T m_2 - w^T m_1)^2}{\sum_{n \in c_1} (w^T x_n - w^T m_1)^2 + \sum_{n \in c_2} (w^T x_n - w^T m_2)^2}$$

$$J(w) = \frac{(w^T m_2 - w^T m_1)(w^T m_2 - w^T m_1)}{\sum_{n \in c_1} (w^T x_n - w^T m_1)(w^T x_n - w^T m_1) + \sum_{n \in c_2} (w^T x_n - w^T m_2)(w^T x_n - w^T m_2)}$$

Since $(w^T m_2 - w^T m_1)$ and $(w^T x_n - w^T m_1)$ are scalars, $aa^T = a^2$

$$J(w) = \frac{(w^T m_2 - w^T m_1)(w^T m_2 - w^T m_1)^T}{\sum_{n \in c_1} (w^T x_n - w^T m_1)(w^T x_n - w^T m_1)^T + \sum_{n \in c_2} (w^T x_n - w^T m_2)(w^T x_n - w^T m_2)^T}$$

$$J(w) = \frac{(w^T m_2 - w^T m_1)(m_2^T w - m_1^T w)}{\sum_{n \in c_1} (w^T x_n - w^T m_1)(x_n^T w - m_1^T w) + \sum_{n \in c_2} (w^T x_n - w^T m_2)(x_n^T w - m_1^T w)}$$

$$J(w) = \frac{w^T (m_2 - m_1)(m_2^T - m_1^T) w}{\sum_{n \in c_1} w^T (x_n - m_1)(x_n^T - m_1^T) w + \sum_{n \in c_2} w^T (x_n - m_2)(x_n^T - m_1^T) w}$$

$$J(w) = \frac{w^T (m_2 - m_1)(m_2 - m_1)^T w}{\sum_{n \in c_1} w^T (x_n - m_1)(x_n - m_1)^T w + \sum_{n \in c_2} w^T (x_n - m_2)(x_n - m_2)^T w}$$

$$J(w) = \frac{w^T S_B w}{w^T S_w w}$$

3. Using the definitions of the between-class and within-class covariance matrices given by $S_B = (m_2 - m_1)(m_2 - m_1)^T$ and $S_w = \sum_{n \in c_1} (x_n - m_1)(x_n - m_1)^T + \sum_{n \in c_2} (x_n - m_2)(x_n - m_2)^T$ respectively, together with $w_0 = -w^T m$ and $m = \frac{1}{N}(N_1 m_1 + N_2 m_2)$ and the choice of target values described in Section 4.1.5, show that the expression

$\sum_{n=1}^N (w^T x_n + w_0 - t_n) x_n = 0$ that minimizes the sum-of-squares error function can be written in the form $(S_W + \frac{N_1 N_2}{N} S_B) w = N(m_1 - m_2)$

$$\sum_{n=1}^N (w^T x_n + w_0 - t_n) x_n = 0$$

$$\sum_{n=1}^N (w^T x_n - w^T m - t_n) x_n = 0$$

$$\sum_{n=1}^N (w^T x_n - w^T \frac{1}{N} (N_1 m_1 + N_2 m_2)) x_n = 0 \text{ since } \sum_{n=1}^N t_n = 0$$

$$\sum_{n=1}^N (w^T x_n - w^T \left(\frac{N_1 m_1}{N} + \frac{N_2 m_2}{N} \right)) x_n = 0$$

$$\sum_{n=1}^N (w^T x_n - w^T \frac{N_1 m_1}{N} - w^T \frac{N_2 m_2}{N}) x_n = 0$$

$$\sum_{n=1}^N (w^T x_n x_n - w^T \frac{N_1 m_1}{N} x_n - w^T \frac{N_2 m_2}{N} x_n) = 0$$

$$\sum_{n=1}^N w^T x_n x_n - \sum_{n=1}^N w^T \frac{N_1 m_1}{N} x_n - \sum_{n=1}^N w^T \frac{N_2 m_2}{N} x_n = 0$$

$$\sum_{n=1}^N w^T x_n x_n = \sum_{n=1}^N w^T \frac{N_1 m_1}{N} x_n + \sum_{n=1}^N w^T \frac{N_2 m_2}{N} x_n$$

$$w^T \sum_{n=1}^N x_n x_n = w^T \frac{N_1 m_1}{N} \sum_{n=1}^N x_n + w^T \frac{N_2 m_2}{N} \sum_{n=1}^N x_n$$

$$w^T \sum_{n=1}^N x_n x_n = w^T N_1 m_1 \left(\frac{N_1 m_1}{N} + \frac{N_2 m_2}{N} \right) + w^T N_2 m_2 \left(\frac{N_1 m_1}{N} + \frac{N_2 m_2}{N} \right)$$

$$w^T \sum_{n=1}^N x_n x_n = w^T N_1 m_1 \left(\frac{N_1 m_1}{N} + \frac{N_2 m_2}{N} \right) + w^T N_2 m_2 \left(\frac{N_1 m_1}{N} + \frac{N_2 m_2}{N} \right)$$

4. Show that the Hessian matrix H for the logistic regression model, given by $H = \phi^T R \phi$, is positive definite. Here R is a diagonal matrix with elements $y_n(1 - y_n)$, and y_n is the output of

the logistic regression model for input vector x_n . Hence show that the error function is a concave function of w and that it has a unique minimum.

$$H = \Phi^T R \Phi$$

$$H = \begin{bmatrix} \phi_{11} & \dots & \phi_{n1} \\ \dots & \dots & \dots \\ \phi_{1d} & \dots & \phi_{nd} \end{bmatrix} \begin{bmatrix} y_n(1-y_n) & \dots & 0 \\ \dots & \dots & \dots \\ 0 & \dots & y_n(1-y_n) \end{bmatrix} \begin{bmatrix} \phi_{11} & \dots & \phi_{1d} \\ \dots & \dots & \dots \\ \phi_{n1} & \dots & \phi_{nd} \end{bmatrix}$$

Since y_n is always a positive value between 0 and 1, $y_n(1-y_n)$ will always be positive

$$H = \begin{bmatrix} \phi_{11}y_n(1-y_n) & \dots & \phi_{n1}y_n(1-y_n) \\ \dots & \dots & \dots \\ \phi_{1d}y_n(1-y_n) & \dots & \phi_{nd}y_n(1-y_n) \end{bmatrix} \begin{bmatrix} \phi_{11} & \dots & \phi_{1d} \\ \dots & \dots & \dots \\ \phi_{n1} & \dots & \phi_{nd} \end{bmatrix}$$

$$H = \begin{bmatrix} \phi_{11}y_n(1-y_n)\phi_{11} + \dots \phi_{n1}y_n(1-y_n)\phi_{n1} & \dots & \phi_{11}y_n(1-y_n)\phi_{1d} + \dots \phi_{n1}y_n(1-y_n)\phi_{n1} \\ \dots & \dots & \dots \\ \phi_{11}y_n(1-y_n)\phi_{11} + \dots \phi_{n1}y_n(1-y_n)\phi_{n1} & \dots & \phi_{11}y_n(1-y_n)\phi_{11} + \dots \phi_{n1}y_n(1-y_n)\phi_{n1} \end{bmatrix}$$

Since we know that $X^T X = X^T I X$ is always positive definite and we have $\Phi^T R \Phi$

where R is a diagonal matrix where all diagonal entries are between 0 and 1, we know

R will act similar to I in that it won't change the positive definite nature of $\Phi^T \Phi$

Linear Regression: Experiment

This program was implemented in Python 3.8, using numpy to handle linear algebra operations, matplotlib.pyplot to handle plotting, scipy.io to open and read from the diabetes.mat file, and math to apply natural logarithms to lambda.

The program starts by extracting data from the diabetes.mat file. It then estimates the optimal weight vector by applying ridge regression. However, to apply ridge regression, you need a lambda value (regularization term) to control the effect of the regularization term. If lambda is too big, you end up underfitting your data; if lambda is too small, you end up overfitting your data. Varying lambdas were tested to estimate the weight vector, from .00001, .0001, ... 1, 10. With these weight vectors, the predicted y training and testing values were generated. Finally, they were compared with the ground truth through using the Mean Squared Error formula.

After generating these different MSEs for the training and testing error, 5-fold-cross-validation was performed on the training data matrix. Through this, an optimal lambda value was generated. Once more, the ridge regression was called to estimate the weight vector, this time using the optimal lambda. These results were added to previously created lists of generated MSEs for the training and testing datasets. Finally, a plot was created to show the MSE varying with the log of lambda. The code can be seen below and is also attached to the submission.

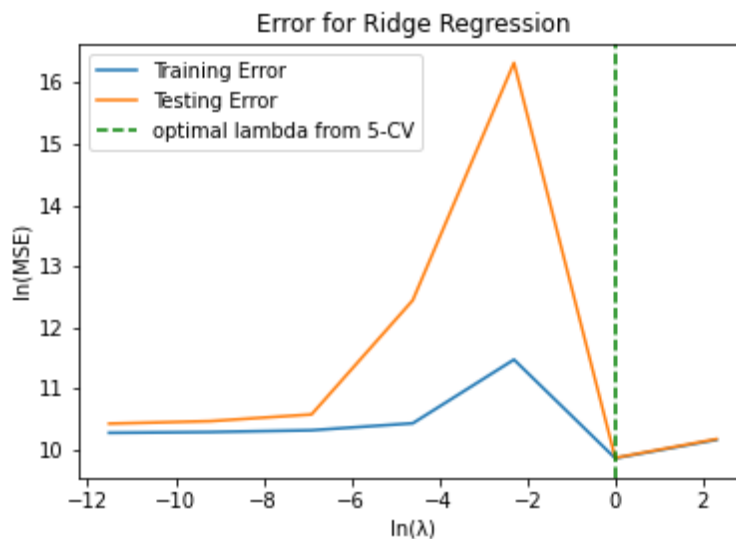


Fig. [1]. A chart showing how the natural log of the mean squared error varies with the natural log of different lambdas, along with the optimal lambda discovered through 5-fold cross validation

Above, a chart can be seen, generated from the program below. This chart shows the varying mean squared error for varying lambdas, again ranging from .00001, .0001, ... 1, 10. The natural logarithm was applied to the mean squared error to increase the distinguishability of the error between different lambdas, as the natural log is a monoatomic increasing function. Based on this experiment, it is discovered that a lambda of 1 ($\ln(\lambda) = 0$) is the optimal lambda, as based on the cross-validation error from using 5 folds. Using a lambda of 0 in ridge regression for the regularization parameter should yield the most accurate results for this diabetes dataset.

```
import matplotlib.pyplot as plt
import scipy.io
import numpy as np
import math

def plot_err(x, y_train, y_test, lam_opt):
    plt.ylabel("ln(MSE)")
    plt.xlabel("ln(λ)")
    plt.title("Error for Ridge Regression")
    plt.plot(x, y_train, label="Training Error")
    plt.plot(x, y_test, label="Testing Error")
    plt.axvline(x=lam_opt, color='g', linestyle='--',
               label="optimal lambda from 5-CV")
    plt.legend(loc="upper left")
    plt.savefig("lamda_err.png")
    plt.show()

def MSE(y_pred, y_truth):
    mse = 0
    for i in range(len(y_pred)):
        mse += ((y_truth[i] - y_pred[i]) ** 2)
```

```

    return math.log(mse / len(y_pred))

def ridge_regression(lam, X, y):
    # Finds the optimal weights using ridge regression i.e. least squares
    # with a regularization parameter defined by the l2-norm
    xtx = np.matmul(np.transpose(X), X)
    xtx_dim = np.shape(xtx)[0]
    w = (xtx + lam * np.identity(xtx_dim))
    w = np.linalg.inv(w)
    w = np.matmul(w, np.transpose(X))
    w = np.dot(w, y)
    return w

def error(k, pred, truth):
    total = 0
    for i in range(len(pred)):
        total += ((pred[i] - truth[i]) ** 2)
    return total

def cross_val(k, x_train, y, lams):
    fold_len = len(x_train) // k
    # Partition data into k-folds
    D, Y = [], []
    for i in range(0, len(x_train), fold_len):
        D.append(x_train[i:i + fold_len])
        Y.append(y[i:i + fold_len])
    D = D[:-1] # Got ride of the array of size 2 because bugs
    Y = Y[:-1]

    # Goes through every lambda and every subset
    all_cv = []
    for lam in lams:
        all_fitted = []
        for i in range(len(D)):
            test = D.pop(i)
            y_test = Y.pop(i)
            train = np.array([j for i in D for j in i])
            y_train = np.array([j for i in Y for j in i])
            D.insert(i, test)
            Y.insert(i, y_test)

            # Train the model with subset
            w = ridge_regression(lam, train, y_train)
            y_pred = np.dot(train, w)
            err = error(k, y_pred, y)
            all_fitted.append(err[0])

        # Compute average error for a specific lambda
        cv_err = (1 / k) * sum(all_fitted)
        all_cv.append((cv_err, lam))

    # return minimum cv error's lamda
    return min(all_cv, key=lambda t: t[0])

mat = scipy.io.loadmat('diabetes.mat')

# Get the data
x_train = np.array(mat["x_train"])
x_test = np.array(mat["x_test"])

```

```

y_train = np.array(mat["y_train"])
y_test = np.array(mat["y_test"])

# Test different lambdas to see which one is best ;)
lams = [.00001, .0001, .001, .01, .1, 1, 10]
lams = [math.log(l) for l in lams]
y_train_err, y_test_err = [], []
for lam in lams:
    # Finds training error
    w_train = ridge_regression(lam, x_train, y_train)
    y_pred_train = np.dot(x_train, w_train)
    err_train = MSE(y_pred_train, y_train)
    y_train_err.append(err_train)

    # Finds testing error
    w_test = ridge_regression(lam, x_test, y_test)
    y_pred_test = np.dot(x_test, w_test)
    err_test = MSE(y_pred_test, y_test)
    y_test_err.append(err_test)

# Compute 5-cross fold validation and add the optimal lambda to results
lam = cross_val(5, x_train, y_train, lams)
plot_err(lams, y_train_err, y_test_err, lam[1])

```